

Bayesian Learning

Computer Lab 4

Elvin Granat (elvgr805)

Max Fischer (maxfi539)

1) Poisson regression, the MCMC way.

a)

To obtain the maximum likelihood estimator of β in the Poisson regression model for the eBay dataset, we initially imported the data from the .dat-file.

```
### Setup
data.ebay = read.table("eBayNumberOfBidderData.dat", header=TRUE)
```

A Poisson regression model was then fitted to the data, with nBids as ground truth and the rest of the data as covariates.

```
# a)
# Fit the model
glm.fit = glm(nBids ~ .-Const, data=data.ebay, family="poisson")

# Check significance
summary(glm.fit)
```

Significance of variables:

| Covariate | Estimate | Std | z-value | Significance |
|-------------|----------|---------|---------|------------------------|
| (Intercept) | 1.07244 | 0.03077 | 34.848 | $\Pr(> z) < 2e-16$ |
| PowerSeller | -0.02054 | 0.03678 | -0.558 | $\Pr(> z) = 0.5765$ |
| VerifyID | -0.39452 | 0.09243 | -4.268 | $\Pr(> z) = 1.97e-05$ |
| Sealed | 0.44384 | 0.05056 | 8.778 | $\Pr(> z) < 2e-16$ |
| MinBlem | -0.05220 | 0.06020 | -0.867 | $\Pr(> z) = 0.3859$ |
| MajBlem | -0.22087 | 0.09144 | -2.416 | $\Pr(> z) = 0.0157$ |
| LargNeg | 0.07067 | 0.05633 | 1.255 | $\Pr(> z) = 0.2096$ |
| LogBook | -0.12068 | 0.02896 | -4.166 | $\Pr(> z) = 3.09e-05$ |
| MinBidShare | -1.89410 | 0.07124 | -26.588 | $\Pr(> z) < 2e-16$ |

From the summary above, we can draw the conclusion that the following covariates' confidence interval do not include zero (with the probability written afterwards) and thus are significant:

- (Intercept) (99,9%)
- VerifyID (99,9%)
- Sealed (99,9%)
- LogBook (99,9%)
- MinBidShare (99,9%)
- MajBlem (95%)

b)

A Bayesian analysis of the Poisson regression was done.

- Zellner g-prior: $\beta \sim N(0, 100 * (X^T X)^{-1})$ (X is the n x p covariate matrix)
- Likelihood: $y_i | \beta \sim \text{Poisson}[\exp(x_i \beta)]$, $i = 1, \dots, n$
- Posterior: $\beta | y_i \sim N[\hat{\beta}, J_y^{-1}(\hat{\beta})]$, $\hat{\beta}$: Mode of β , $J_y^{-1}(\hat{\beta})$: Inverse negative Hessian at the mode of β .

To calculate the mode of β and the inverse negative Hessian at the mode of β a function calculating the log posterior was coded:

```
logPostPoisson = function(beta, mu0, Sigma0, X, y) {
  p = length(beta)

  # log of the likelihood
  log.likelihood = sum(y * (X %*% beta) - exp(X %*% beta))

  # if likelihood is very large or very small, steer optim away
  if (abs(log.likelihood) == Inf) log.likelihood = -20000;

  # log of the prior
  log.prior = dmvnorm(beta, mean = mu0, sigma = Sigma0, log = TRUE)

  return(log.likelihood + log.prior)
}
```

The likelihood function was derived in the following way:

$$L(y | \beta) = \prod_{i=1}^n p(y_i | \beta) = \prod_{i=1}^n \frac{(e^{x_i \beta})^{y_i}}{y_i!} e^{-(e^{x_i \beta})} = \prod_{i=1}^n \frac{e^{y_i x_i \beta}}{y_i!} e^{-(e^{x_i \beta})}$$

By using the likelihood function above, the log-likelihood function was derived:

$$\begin{aligned} l(y | \beta) &= \log\left[\prod_{i=1}^n \frac{e^{y_i x_i \beta}}{y_i!} e^{-(e^{x_i \beta})}\right] = \log\left[\prod_{i=1}^n e^{y_i x_i \beta}\right] + \log\left[\prod_{i=1}^n e^{-(e^{x_i \beta})}\right] - \log\left[\prod_{i=1}^n y_i!\right] = \\ &= \log\left[e^{\sum_{i=1}^n y_i x_i \beta}\right] + \log\left[e^{\sum_{i=1}^n -(e^{x_i \beta})}\right] - \log\left[\prod_{i=1}^n y_i!\right] \propto \sum_{i=1}^n y_i x_i \beta * \log[e] + \sum_{i=1}^n -(e^{x_i \beta}) * \log[e] \\ &= \sum_{i=1}^n (y_i x_i \beta - e^{x_i \beta}) \end{aligned}$$

The log-posterior function was numerically optimized by using the optim.R package:

```
y.ebay = data.ebay[,1]
X.ebay = as.matrix(data.ebay[,,-1])
# Prior
beta0 = rep(0, 9) # mean prior
Sigma0 = 100 * solve(t(X.ebay) %*% X.ebay) # Sigma prior

# Optimize posterior of beta given priors and data
optim.res = optim(beta0, logPostPoisson, gr=NULL,
                  beta0, Sigma0, X.ebay, y.ebay,
                  method="BFGS", control=list(fnscale=-1), hessian=TRUE)

# Results
beta.mode = optim.res$par
beta.invhessian = -solve(optim.res$hessian)
```

Mode of beta:

| | β_1 | β_2 | β_3 | β_4 | β_5 | β_6 | β_7 | β_8 | β_9 |
|-------------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| Mode | 1.069841 | -0.020512 | -0.393006 | 0.443556 | -0.052466 | -0.221238 | 0.070697 | -0.120218 | -1.891985 |

Inverse negative Hessian:

| | β_1 | β_2 | β_3 | β_4 | β_5 | β_6 | β_7 | β_8 | β_9 |
|--------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|
| β_1 | 9.45462 5e-04 | -7.138972 e-04 | -2.741517 e-04 | -2.709016 e-04 | -4.454554 e-04 | -2.772239 e-04 | -5.128351 e-04 | 6.436765 e-05 | 1.109935 e-03 |
| β_2 | -7.1389 72e-04 | 1.353076 e-03 | 4.024623 e-05 | -2.948968 e-04 | 1.142960 e-04 | -2.082668 e-04 | 2.801777 e-04 | 1.181852 e-04 | -5.685706 e-04 |
| β_3 | -2.7415 17e-04 | 4.024623 e-05 | 8.515360 e-03 | -7.824886 e-04 | -1.013613 e-04 | 2.282539 e-04 | 3.313568 e-04 | -3.191869 e-04 | -4.292827 e-04 |
| β_{41} | -2.7090 16e-04 | -2.948968 e-04 | -7.824886 e-04 | 2.557778 e-03 | 3.577158 e-04 | 4.532308 e-04 | 3.376467 e-04 | -1.311025 e-04 | -5.759169 e-05 |
| β_5 | -4.4545 54e-04 | 1.142960 e-04 | -1.013613 e-04 | 3.577158 e-04 | 3.624606 e-03 | 3.492353 e-04 | 5.844006 e-05 | 5.854011 e-05 | -6.437066 e-05 |
| β_6 | -2.7722 39e-04 | -2.082668 e-04 | 2.282539 e-04 | 4.532308 e-04 | 3.492353 e-04 | 8.365059 e-03 | 4.048644 e-04 | -8.975843 e-05 | 2.622264 e-04 |
| β_7 | -5.1283 51e-04 | 2.801777 e-04 | 3.313568 e-04 | 3.376467 e-04 | 5.844006 e-05 | 4.048644 e-04 | 3.175060 e-03 | -2.541751 e-04 | -1.063169 e-04 |
| β_8 | 6.43676 5e-05 | 1.181852 e-04 | -3.191869 e-04 | -1.311025 e-04 | 5.854011 e-05 | -8.975843 e-05 | -2.541751 e-04 | 8.384703 e-04 | 1.037428 e-03 |
| β_9 | 1.10993 5e-03 | -5.685706 e-04 | -4.292827 e-04 | -5.759169 e-05 | -6.437066 e-05 | 2.622264 e-04 | -1.063169 e-04 | 1.037428 e-03 | 5.054757 e-03 |

c) We now simulate from the actual posterior of β using Metropolis algorithm. To do this we program a general function that uses the Metropolis algorithm from any posterior distribution.

The proposal density, given the usage of random walk Metropolis is:

$$\theta_p | \theta_c \sim N(\theta_c, \tilde{c} \cdot \Sigma)$$

Where:

$$\Sigma = J_y^{-1}(\tilde{\beta})$$

(Inverse negative hessian from b))

And c is a given tuning parameter.

```
Metropolis = function(nBurnIn, nSamples, theta, c, logPostFunc, ...) {
```

The Metropolis function takes in a number of parameters:

nBurnIn = number of burnin iterations of the algorithm

nSamples = number of samples after the burnin iterations

theta = the parameter of which to sample, used to evaluate the posterior density

c = tuning parameter for the proposal density

logPostFunc = computes the log posterior density of the parameters

... = prior hyperparameters used in logPostFunc together with theta to calculate the log posterior density

Final Metropolis function is as follows:

```
Metropolis = function(nBurnIn, nSamples, theta, c, logPostFunc, ...) {  
  # Setup  
  theta.c = theta  
  Sigma.c = c *  $\beta$ .invhessian  
  nAccepted = 0  
  p = length(theta)  
  theta.samples = matrix(NA, nSamples, p)  
  
  # Iterations  
  for(i in -nBurnIn : nSamples) {  
    # Sample from proposal distribution  
    theta.p = as.vector(rmvnorm(1, mean = theta.c, sigma = Sigma.c))  
  
    # Calculate log posteriors  
    log.post.p = logPostFunc(theta.p, ...)  
    log.post.c = logPostFunc(theta.c, ...)  
  
    # Calculate alpha  
    alpha = min(1, exp(log.post.p - log.post.c))  
  
    # Select sample with probability alpha  
    u = runif(1)  
    if (u <= alpha){  
      theta.c = theta.p  
      nAccepted = nAccepted + 1  
    }  
  
    # Save sample if not burnin  
    if (i>0) theta.samples[i,] = theta.c  
  }  
  cat("Sampled", nSamples, "samples with an acceptance rate of", nAccepted/nSamples)  
  return(theta.samples)  
}
```

Using number of burnin iterations, number of samples and tuning parameter c of:

```
# Setup sampling  
c = 0.5  
nSamples = 4000  
nBurnIn = 1000
```

And calling the function as follows:

```
# Samples from posterior using metropolis  
 $\beta$ .samples = Metropolis(nBurnIn, nSamples,  $\beta$ .mode, c, logPostPoisson,  $\beta_0$ , Sigma0, X.ebay, y.ebay)
```

Results:


```
# Estimate parameters using posterior mean  
β.post.mean = apply(β.samples, 2, mean)
```

Mean of beta:

| | β_1 | β_2 | β_3 | β_4 | β_5 | β_6 | β_7 | β_8 | β_9 |
|------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| Mean | 1.075024 | -0.027237 | -0.405541 | 0.444659 | -0.056203 | -0.234342 | 0.073121 | -0.119552 | -1.888617 |

To graphically compare the parameter values of the different estimation methods we plot the value difference of the parameters. With the result of glm method as the middle values.


Results in the following plot:

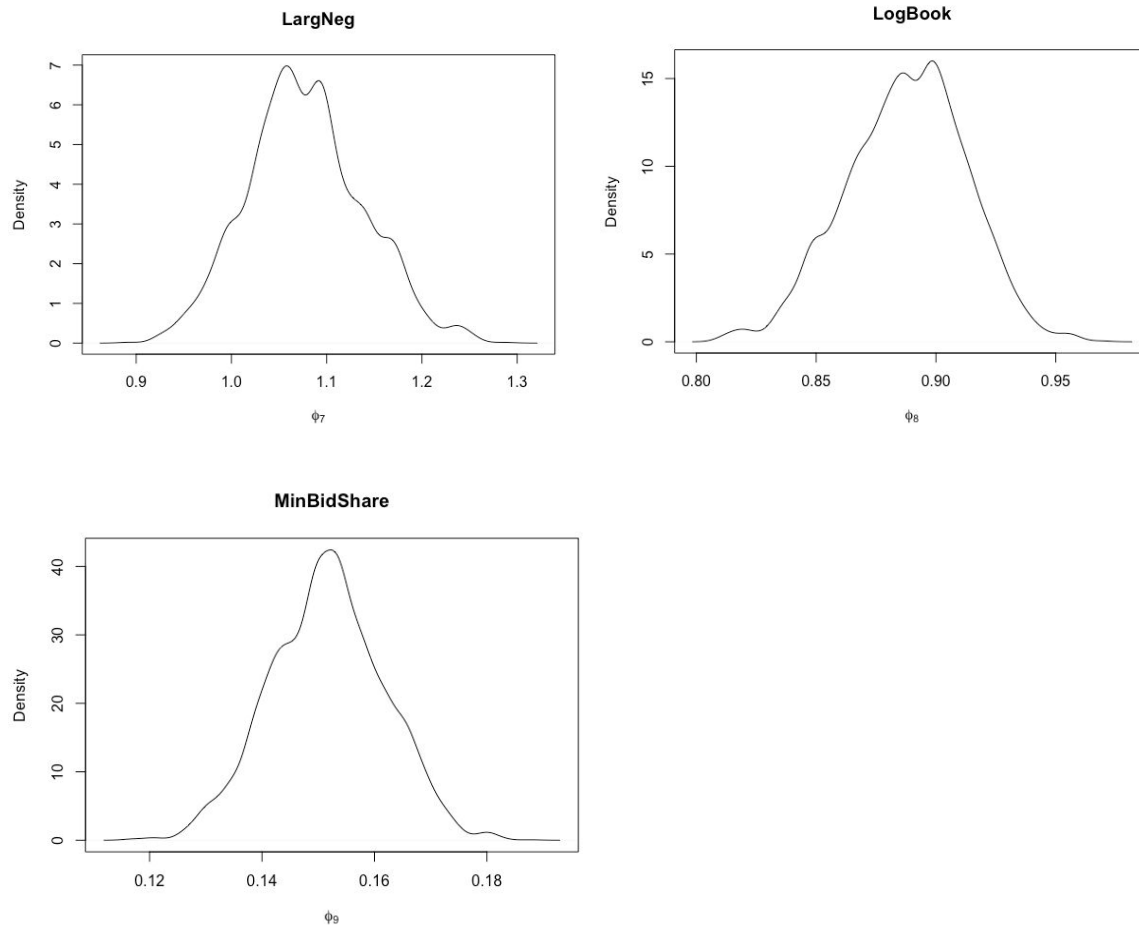


We can see that glm and posterior mode using optim results in similar results whereas MCMC gives very different parameter values for some parameters.

To visualize the convergence MCMC we compute the posterior distribution of:

$\phi_j = \exp(\beta_j)$ for the sampled betas.





Which looks like reasonable posterior distributions. Interesting that we can see that the range of the parameter values that were deemed significant in a) don't cross 1 with any of the samples. Whereas the not so significant covariates' parameters have most of its mass close to 1.

d)

We now use the draws from c) to calculate the probability that a certain auction will have 0 bids.

Given auction:

- PowerSeller = 1
- VerifyID = 1
- Sealed = 1
- MinBlem = 0
- MajBlem = 0
- LargNeg = 0
- LogBook = 1
- MinBidShare = 0.5

Where we calculate the probability using samples from the poisson distribution with lambda value calculated using the samples from the MCMC. And finally calculating the probability of having 0 bids using the proportion of the samples that had 0 bids.

```
# Auction vector
x = c(1, 1, 1, 1, 0, 0, 0, 1, 0.5)


# Setup for sampling
nBids.samples = numeric()

# Sample nBids
for(i in 1:nSamples) {
  nBids.samples[i] = rpois(1, exp(x %*% β.samples[i,]))
}

# Plot distribution
barplot(table(nBids.samples))

# Calculate probability of no bidders
prob = sum(nBids.samples == 0) / nSamples # 0.36
```

With corresponding distribution:



Resulting in a **36%** probability of 0 bids.